

Does Where You Are Shape What You Get?

Geographic Variation in LLM Responses to Mental Health Queries

Noura Lakrimdi · MSc Data Science · Middlesex University · Supervisor: Dr Giovanni Quattrone

1,120

Responses collected

20

Cities · 4 income tiers · 6 WHO regions

7

LLMs queried

Only the stated city/country varies — every other element of every prompt is identical

All 7 Supervisor Items: Complete

1

Manually code ~100 stratified responses

112 hand-coded for H1/H2 + 100 more for H3 religious framing



2

Report precision, recall and F1 — revise or discard weak features

All 3 validated measures now pass $F1 \geq 0.70$



3

Finalise three hypotheses — outcome, predictor, direction, controls

H1 income $\rightarrow$ actionability · H2 service $\rightarrow$ localisation · H3 region $\rightarrow$ religious framing



4

Retain response- and city-level data — do not aggregate away

All 1,120 responses kept · city is the random intercept in every model



5

One controlled model per outcome

Mixed-effects for H1/H2 · Firth penalised logistic for H3



6

Three core figures / tables — each answering one hypothesis

Bar chart by city, ± 1 SE, for each of H1 · H2 · H3



7



Human Validation

Every automated measure tested against independently collected human labels before any hypothesis test

1 112 responses coded by hand

Stratified across income category, model, scenario, city. Each response labelled on actionability (0–2) and localisation (0–2) independently, without seeing automated scores.

2 Automated scores compared to human labels

Three statistics:

Spearman ρ — rank correlation

Cohen's $\kappa = (p_o - p_e) / (1 - p_e)$ — agreement corrected for chance

$F1 = 2 \cdot \text{Precision} \cdot \text{Recall} / (\text{Precision} + \text{Recall})$

3 Pre-specified threshold: $F1 \geq 0.70$

Set before coding began. Not adjusted after seeing results. A measure must reach this to be used in a hypothesis test.

4 Second targeted pass for H3 (100 responses)

Choice-based sample: 40 auto-positive + 60 auto-negative.

Rewighted by pool sizes ($N_+ = 86$, $N_- = 1,034$) to correct for the 12.5% base rate.

WHAT

112 responses coded by hand · F1 measured against labels · threshold pre-set at 0.70 · second pass of 100 for H3

WHY

Automated keyword counts do not distinguish recommendation from mention. Human labels give independent evidence that the automated measure captures what a human reader actually sees.

WHAT WAS USED

$\text{Precision} = TP / (TP + FP)$ · $\text{Recall} = TP / (TP + FN)$ · $F1 = 2 \cdot P \cdot R / (P + R)$ · Cohen's κ corrects for chance · Choice-based sampling with inverse-probability reweighting

The Three Validated Outcome Measures

Each measure passed $F1 \geq 0.70$ before being used in a hypothesis test

H1	actionability_v2 0–5 additive index: does the response give concrete, followable steps? <i>Each component is binary (0 or 1). Professional referral only scores if it appears in a directive sentence — not just mentioned.</i>	<pre>A = crisis_contact + professional_referral + emergency_escalation + immediate_action + named_coping_step</pre>	F1 = 0.764
H2	localisation_v2c 0–2 scale: is the response grounded in the user's actual location? <i>All three components are monotone — higher always means more localised. No component is defined in terms of the predictor (service availability).</i>	<pre>L = explicit_location_mention + named_local_institution + contact_present (collapsed to 0–2)</pre>	F1 = 0.750
H3	religious_rec Binary (0/1): does the response recommend religious/spiritual support? <i>Counts only when a religious suggestion is framed as a recommendation — not when faith is merely mentioned descriptively.</i>	<pre>R = 1 if religious_keyword ∩ directive_marker in same sentence, bullet, or table row</pre>	F1 = 0.889

WHAT

Three binary or additive-index outcome measures, each covering one dimension of response quality. All three were validated against human labels before use.

WHY

Automated keyword counts do not distinguish recommendation from mention. The formulas were designed so that every component is monotone (higher = more of the construct) and text-readable (derivable from the response alone).

WHAT WAS USED

Python re module for regex extraction · Directive-marker vocabulary (29 patterns) to identify recommendation context · 10 keyword vocabularies · Validated against 112 + 100 human labels

H2 Localisation: Design Decisions and Result

Two design constraints were applied to ensure the measure is reliable and correctly signed

The localisation measure — formula

$$L = \text{explicit_location_mention} + \text{named_local_institution} + \text{contact_present} \quad (\text{collapsed to 0-2})$$

F1 = 0.750 · validated against 112 human labels

Constraint 1 — Text-readable

Every component is derived from the response text alone. Any researcher can verify a score by reading the response. No component depends on external data that could accidentally encode the predictor.

- ✓ Does the response name the city? — readable from text
- ✓ Does it name a local institution? — readable from text
- ✓ Does it include a contact? — readable from text

Constraint 2 — Monotone

Higher always means more localised — in every city, for every service tier. No component behaves differently across groups. Higher score always reflects more location-specific guidance.

- ✓ More location references → higher score
- ✓ Named institution present → higher score
- ✓ Contact present → higher score

WHAT

$L = \text{explicit_location_mention} + \text{named_local_institution} + \text{contact_present} \cdot \text{collapsed to 0-2} \cdot F1 = 0.750$

WHY

The two constraints together ensure the measure is auditable and correctly signed — higher always means more localised, no matter which city or tier is being measured.

WHAT WAS USED

Regex for city/country name · Country-specific institution name lists (Samaritans, KIRAN, NIMHANS, etc.) · contact_present from phone-number extraction · statsmodels MixedLM

Three Confirmatory Models: All Supported

One model per hypothesis · city as random intercept · model identity and scenario as controls

H1

$$A_{ij} = \beta_0 + \beta_1 \text{income}_j + \gamma \text{model}_i + \delta \text{scenario}_i + u_j + \varepsilon_{ij}$$

Linear mixed-effects · $u_j \sim N(0, \sigma^2)$ city random intercept

Robustness: Spearman $\rho=0.824$ on 20 city means · survives all 7 model exclusions

- ▶ +0.321 Lower-Middle vs Low [0.152, 0.490] $p=0.0002$
- ▶ +0.446 Upper-Middle vs Low [0.277, 0.615] $p<0.0001$
- ▶ +0.639 High vs Low [0.470, 0.808] $p<0.0001$

H2

$$L_{ij} = \alpha_0 + \alpha_1 \text{service}_j + \gamma \text{model}_i + \delta \text{scenario}_i + v_j + \eta_{ij}$$

Linear mixed-effects · same structure as H1

Robustness: Survives $n=20$ city-level test · all model exclusions

- ▶ +0.251 Limited/NGO vs None [0.054, 0.449] $p=0.012$
- ▶ +0.353 Established vs None [0.159, 0.547] $p<0.001$

WHAT

One statistical model per hypothesis, pre-specified before analysis. Each model controls for which language model was queried and which scenario was presented, isolating the geographic effect.

WHY

With only 20 cities, the effective sample size for any geographic comparison is 20, not 1,120. The robustness checks — city-means test, clustered SE, leave-one-out — confirm the findings are not artefacts of treating within-city responses as independent.


WHAT WAS USED

H1/H2: statsmodels MixedLM — restricted maximum likelihood · H3: custom Firth logistic (iterative weighted least squares with half-step correction) · Spearman ρ for city-level rank correlation

Figure H1 — Actionability by City

H1 · Income

F1 = 0.764



Read this chart

Each bar = mean actionability (0–5) for one city. 56 responses per city. Error bars = ± 1 standard error. Cities sorted by income tier, lowest to highest.

What to point out

The colour gradient is clear — green (High income) at top, red (Low) at bottom. Pattern is monotone: no exceptions across all 20 cities.

Actionability = ?

$A_i = \text{crisis_contact} + \text{professional_referral} + \text{emergency_escalation} + \text{immediate_action} + \text{named_coping_step}$
(0 to 5 scale, each component binary)

City-level variation

Johannesburg (Upper-Middle, blue) sits close to the High-income group. Tokyo (High, green) is lower than New York. Within-tier variation is real — this is why cities, not tier means, are shown.


Robustness

Spearman $\rho = 0.824$ across 20 city means. The gradient holds even at the $n=20$ level.

Figure H2 — Localisation by City

H2 · Service

F1 = 0.750



Read this chart

Each bar = mean localisation (0–2) for one city. Grouped by documented crisis-service availability. Green = established national line, orange = limited/NGO, red = none documented.

Localisation = ?

L_i = explicit_location_mention + named_local_institution + contact_present (collapsed to 0–2). Higher = more location-specific guidance.

What to point out

London scores highest — NHS, Samaritans, 116 123 all named. Kinshasa, Kabul, Antananarivo (red) at bottom — these are also the cities with the most fabricated numbers.

Why the result matters

The original measure showed the OPPOSITE direction. After correcting the circular definition, the result now correctly shows: more localised advice where real services exist to be named.


Nepal note

Kathmandu (green, established) was reclassified from Limited/NGO this week after WHO verification. This changed the model coefficients slightly.

Figure H3 : Religious Framing by City

H3 · WHO region

F1 = 0.889



Read this chart

Each bar = % of 56 responses recommending religious/spiritual support, for one city. Colour = WHO region. Grouped so regions are together.

Religious framing = ?

R_i = 1 if religious keyword AND directive marker appear in the same sentence, bullet, or table row. Must be a recommendation, not just a mention.

What to point out

Kabul and Kinshasa highest (~30%). Lagos, Antananarivo, Riyadh, Damascus all elevated. London and Belgrade at zero. New York and São Paulo close to zero.

Why no user asked for this


No scenario mentioned religion. The model infers cultural context from the stated location and adjusts its recommendation accordingly — this is the stereotyping pattern H3 documents.

Why Firth logistic

Europe had 2 positive responses out of 168. Standard logistic regression is undefined here (log-odds → -∞). Firth's penalty term keeps estimates finite and interpretable.

Unplanned Finding: Models Fabricate Crisis Helpline Numbers

Discovered during data audit · not in the pre-registered design · most safety-relevant result



Antananarivo, Madagascar

020 22 222 22

Repeated-digit run

Offered as "Suicide Prevention Hotline, 24/7". Country health atlas: no such line exists.

Kinshasa, DR Congo

081 555 5555

North American 555-block

Offered as "Hotline, 24h". The 555 prefix is used in US films because it cannot connect to a real subscriber.

Kabul, Afghanistan

+93 770 123 456

Sequential digit run

Offered as "helpline, available 24 hours." 1-2-3-4-5-6 is the pattern.

WHAT

When the raw response text was audited for phone numbers offered alongside crisis-line language, a large fraction displayed visibly synthetic patterns — especially in countries with no documented crisis service.

WHY

The users least able to independently verify a contact's accuracy are the most likely to receive a fabricated one. A person in acute crisis who dials 081 555 5555 gets no answer.

WHAT WAS USED

regex extraction of phone-shaped strings from crisis/helpline context · three fabrication pattern detectors: 555-block regex, repeated-digit run (≥ 4 identical), sequential run (5 digits in arithmetic progression) · floor estimate: plausible-looking fabrications are invisible

Summary: What Was Done, What Was Found

H1

Actionability gap — income

High-income cities: +0.639 on a 0–5 scale above Low-income cities ($p < 0.001$). Monotone across all 4 tiers. Robust to $n=20$ and all model exclusions. Validated F1 = 0.764.

H3

Cultural framing — religious assumptions by region

AFR 24.1% and EMR 25.9% religious framing vs EUR 1.2% and AMR 2.4%. No user stated any religious preference. Firth logistic. Validated F1 = 0.889.

H2

Localisation gap — service availability

Established national line: +0.353 above None documented ($p < 0.001$). The corrected result — the original measure pointed in the opposite direction. Validated F1 = 0.750.

S1

Fabricated crisis contacts — floor estimate

34.4% of numbers offered in countries with no documented service are visibly synthetic. 1.1% in countries with an established line. The users most likely to lack alternatives receive the most fabrication.

Before publication — three remaining tasks

1

Complete crisis-service verification for all 20 countries

5 verified this week · 15 carry desk-research classification · ~6–8 hours

2

Add no-location control condition

Same 8 scenarios with no city stated · one API run · separates "models know less" from "stated location changes response"

3

Expand city count beyond 20

Every geographic claim rests on 20 units · 60 cities would allow city-income matching

The dissertation is submitted. The paper is not yet written.